

F1 Prediction

FORMULA 1 CHAMPIONSHIP EDA AND WINNER PREDICTION

Loading libraries

```
library(ggplot2)
library(dplyr)
library(rpart)
library(corrplot)
library(treemap)
library(treemapify)
library(tree)
library(randomForest)
library(caret)
library(e1071)
library(rpart.plot)
library(car)
```

Data loading

```
df = read.csv("~/Library/Mobile Documents/com~apple~CloudDocs/UNICATT/Data analysis techniques and tools")
str(df)
```

```
## 'data.frame': 23693 obs. of 25 variables:
## $ raceId      : int 1 1 1 1 1 1 1 1 1 ...
## $ driverId    : int 10 15 16 17 18 2 20 21 22 3 ...
## $ constructorId: int 7 7 10 9 23 2 9 10 23 3 ...
## $ circuitId   : int 1 1 1 1 1 1 1 1 1 ...
## $ resultId    : int 7557 7556 7562 7565 7554 7563 7566 7564 7555 7559 ...
## $ number       : int 10 9 20 14 22 6 15 21 23 16 ...
## $ grid         : int 19 20 16 8 1 9 3 15 2 5 ...
## $ positionOrder: int 4 3 9 12 1 10 13 11 2 6 ...
## $ points       : num 5 6 0 0 10 0 0 0 8 3 ...
## $ laps          : int 58 58 58 57 58 58 56 58 58 58 ...
## $ fastestLapSpeed: num 216 215 215 216 217 ...
## $ status        : chr "Finished" "Finished" "Finished" "+1 Lap" ...
## $ dob           : chr "1982-03-18" "1974-07-13" "1983-01-11" "1976-08-27" ...
## $ driv_nationality: chr "German" "Italian" "German" "Australian" ...
## $ fullname      : chr "Timo Glock" "Jarno Trulli" "Adrian Sutil" "Mark Webber" ...
## $ round         : int 1 1 1 1 1 1 1 1 1 ...
## $ date          : chr "2009-03-29" "2009-03-29" "2009-03-29" "2009-03-29" ...
```

```

## $ const_name      : chr  "Toyota" "Toyota" "Force India" "Red Bull" ...
## $ name           : chr  "Albert Park Grand Prix Circuit" "Albert Park Grand Prix Circuit" "Albert ...
## $ const_points    : num  11 11 0 0 18 0 0 0 18 3 ...
## $ const_wins      : int  0 0 0 0 1 0 0 0 1 0 ...
## $ driver_age     : int  27 35 26 33 29 32 22 36 37 24 ...
## $ fastestLap_ms  : num  88416 88916 88943 88508 88020 ...
## $ winner         : int  0 0 0 0 1 0 0 0 0 0 ...
## $ wins           : int  1 1 1 1 2 1 1 1 1 1 ...

head(df)

##   raceId driverId constructorId circuitId resultId number grid positionOrder
## 1      1        10             7       1    7557    10    19          4
## 2      1        15             7       1    7556     9    20          3
## 3      1        16            10      1    7562    20    16          9
## 4      1        17             9       1    7565    14    8           12
## 5      1        18            23      1    7554    22    1           1
## 6      1        2              2       1    7563     6    9           10
##   points laps fastestLapSpeed status      dob driv_nationality
## 1      5   58      215.920 Finished 1982-03-18      German
## 2      6   58      214.706 Finished 1974-07-13      Italian
## 3      0   58      214.640 Finished 1983-01-11      German
## 4      0   57      215.695 +1 Lap 1976-08-27 Australian
## 5     10   58      216.891 Finished 1980-01-19      British
## 6      0   58      216.245 Finished 1977-05-10      German
##   fullname round      date const_name          name
## 1 Timo Glock     1 2009-03-29 Toyota Albert Park Grand Prix Circuit
## 2 Jarno Trulli    1 2009-03-29 Toyota Albert Park Grand Prix Circuit
## 3 Adrian Sutil    1 2009-03-29 Force India Albert Park Grand Prix Circuit
## 4 Mark Webber     1 2009-03-29 Red Bull Albert Park Grand Prix Circuit
## 5 Jenson Button    1 2009-03-29 Brawn Albert Park Grand Prix Circuit
## 6 Nick Heidfeld    1 2009-03-29 BMW Sauber Albert Park Grand Prix Circuit
##   const_points const_wins driver_age fastestLap_ms winner wins
## 1          11        0       27      88416     0    1
## 2          11        0       35      88916     0    1
## 3          0        0       26      88943     0    1
## 4          0        0       33      88508     0    1
## 5         18        1       29      88020     1    2
## 6          0        0       32      88283     0    1

```

Our starting dataframe `df` is composed by 25 columns of features that define for each row a Driver performance in a specific Race of a specific year. So for each row:

- 5 columns describe the unique observation: `raceId`, `driverId`, `constructorId`, `circuitId`, `resultId`. They were useful in the initial part of the data processing to merge multiple features from other datasets into this one main dataframe that include all the features that we later use for data analysis and prediction.
- 2 columns define the performance of the driver in the specific race, mainly `positionOrder` and the binary `winner`. These are the target variables of the regression and classification methods.
- 11 columns are numerical features used:
 - `number`, number of each driver's car

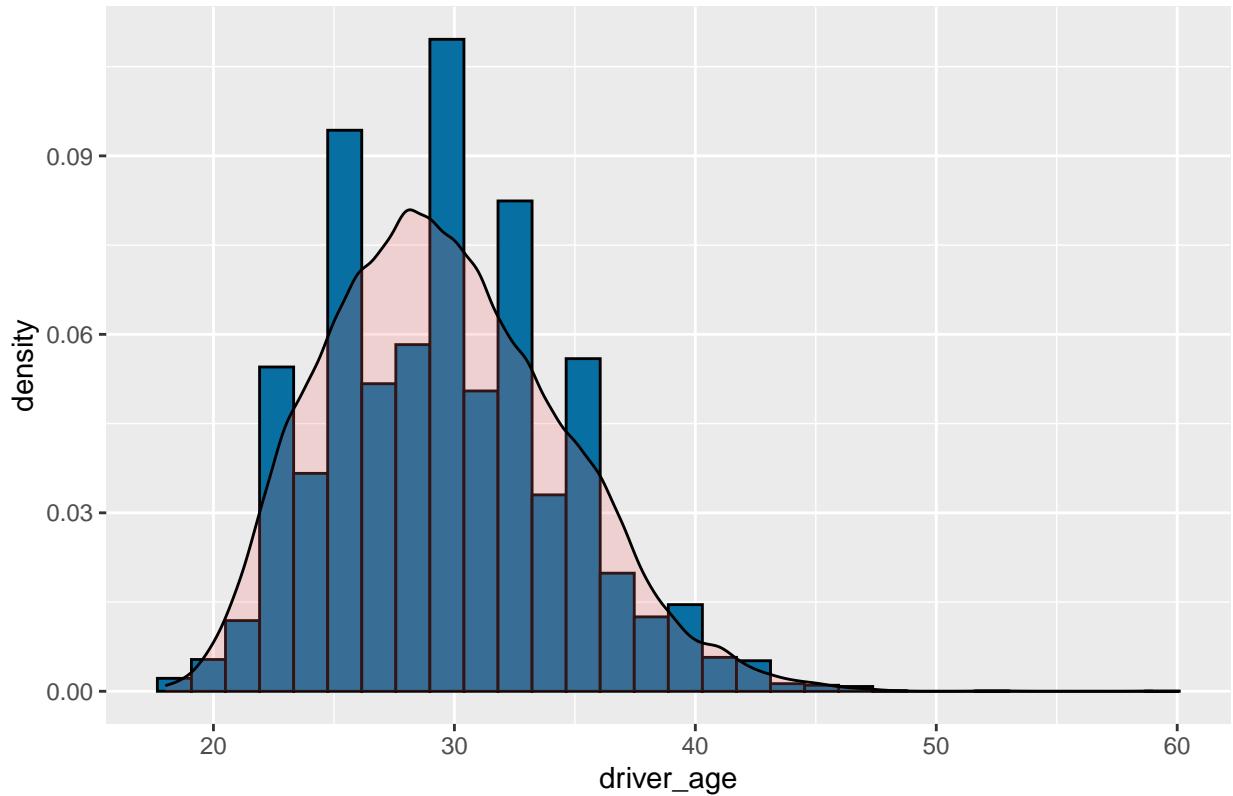
- `grid`, starting position for the driver
 - `points`, number of points earned in the race
 - `laps`, number of laps in the race
 - `fastestLapSpeed`
 - `round`, number of the Race in the season
 - `const_points`, number of points earned by the team before this race
 - `const_wins`, number of wins of the team before the race
 - `driver_age`, age of each driver
 - `fastestLap_ms`, fastest lap in the qualification race
 - `wins`, number of wins before the race for the driver
- 5 columns that describe categorical features:
 - `status`, what happened in the race for the driver
 - `driv_nationality`
 - `fullname`, name of the driver
 - `const_name`, name of the team
 - `name`, name of the circuit in which the race takes place
 - 2 columns for dates: `dob` is birth date for each driver, `date` is the date of the Race.

EDA and plots.

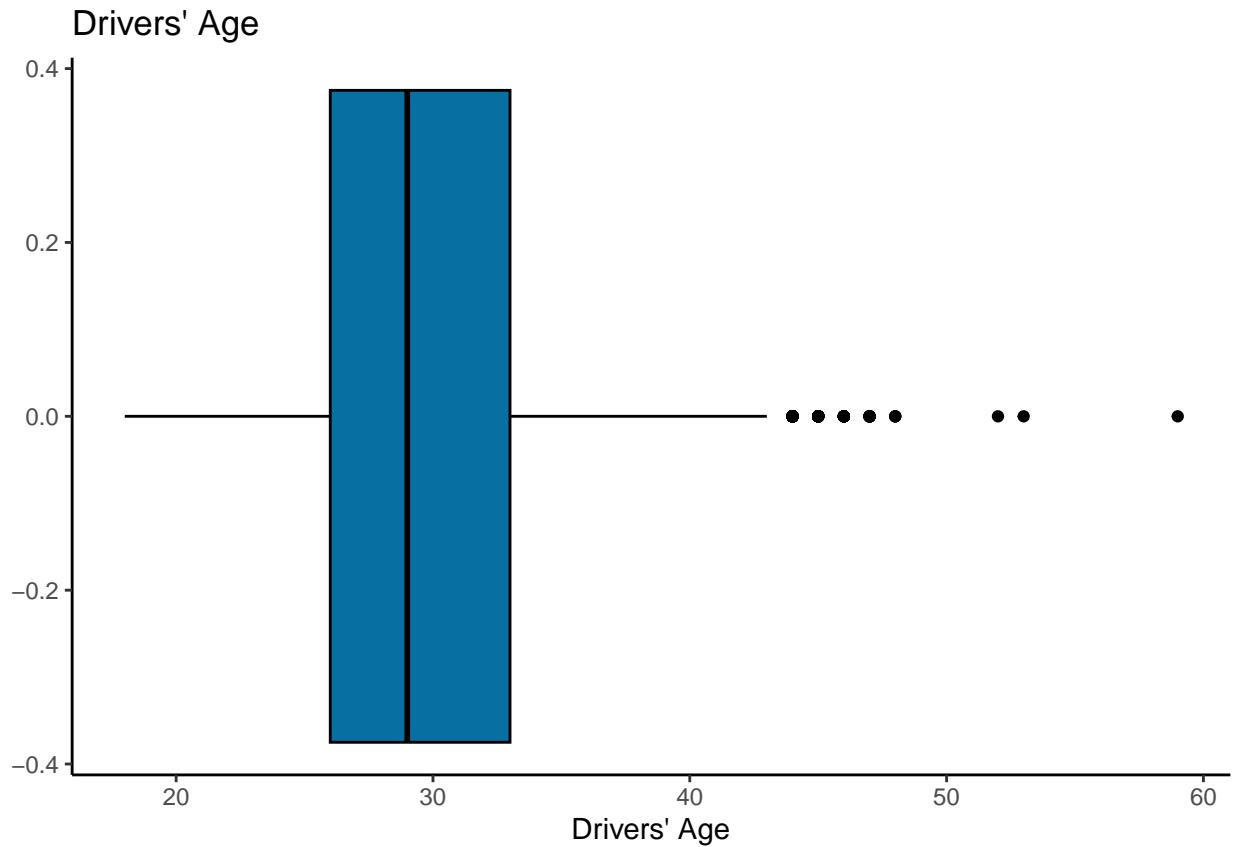
```
ggplot(df, aes(x=driver_age)) +
  geom_histogram(aes(y=after_stat(density)), colour="black", fill= '#076fa2')+
  geom_density(alpha=.2, fill="#FF6666") +
  labs(title = 'Histogram of driver age')

## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
```

Histogram of driver age



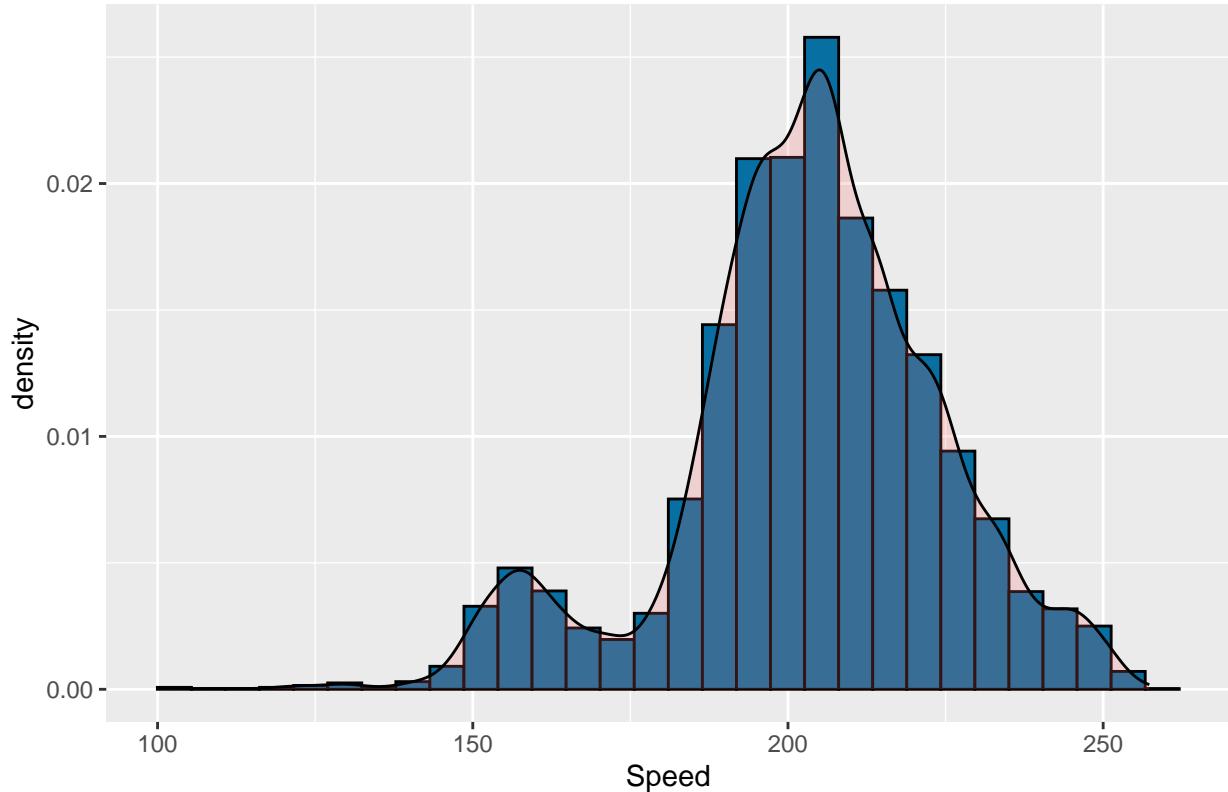
```
ggplot(df, aes(x=driver_age)) +  
  geom_boxplot(fill="#076fa2", color="black") +  
  theme_classic() +  
  labs(title = "Drivers' Age", x = "Drivers' Age")
```



```
ggplot(subset(df, df$fastestLapSpeed > 100), aes(x=fastestLapSpeed)) +  
  geom_histogram(aes(y=after_stat(density)), colour="black", fill= '#076fa2')+  
  geom_density(alpha=.2, fill="#FF6666") +  
  labs(title = 'Histogram for fastest speed in a lap', x = 'Speed')
```

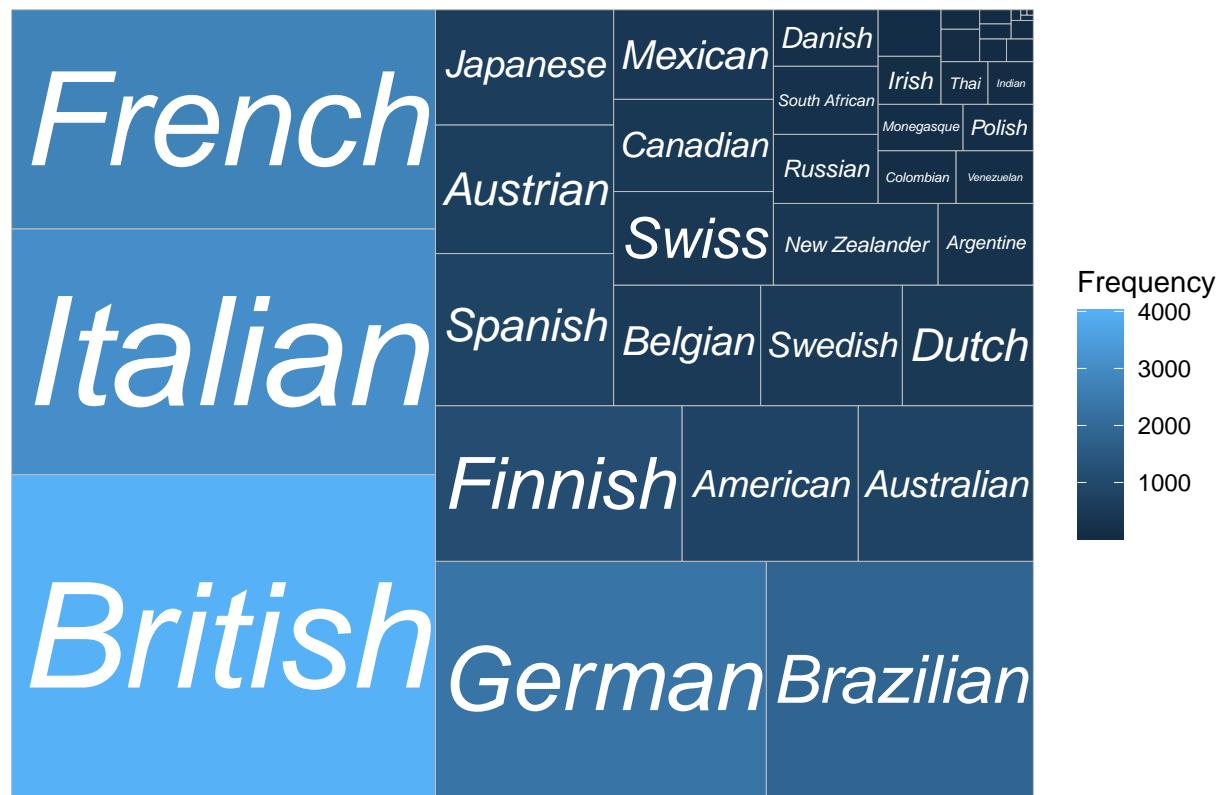
```
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
```

Histogram for fastest speed in a lap



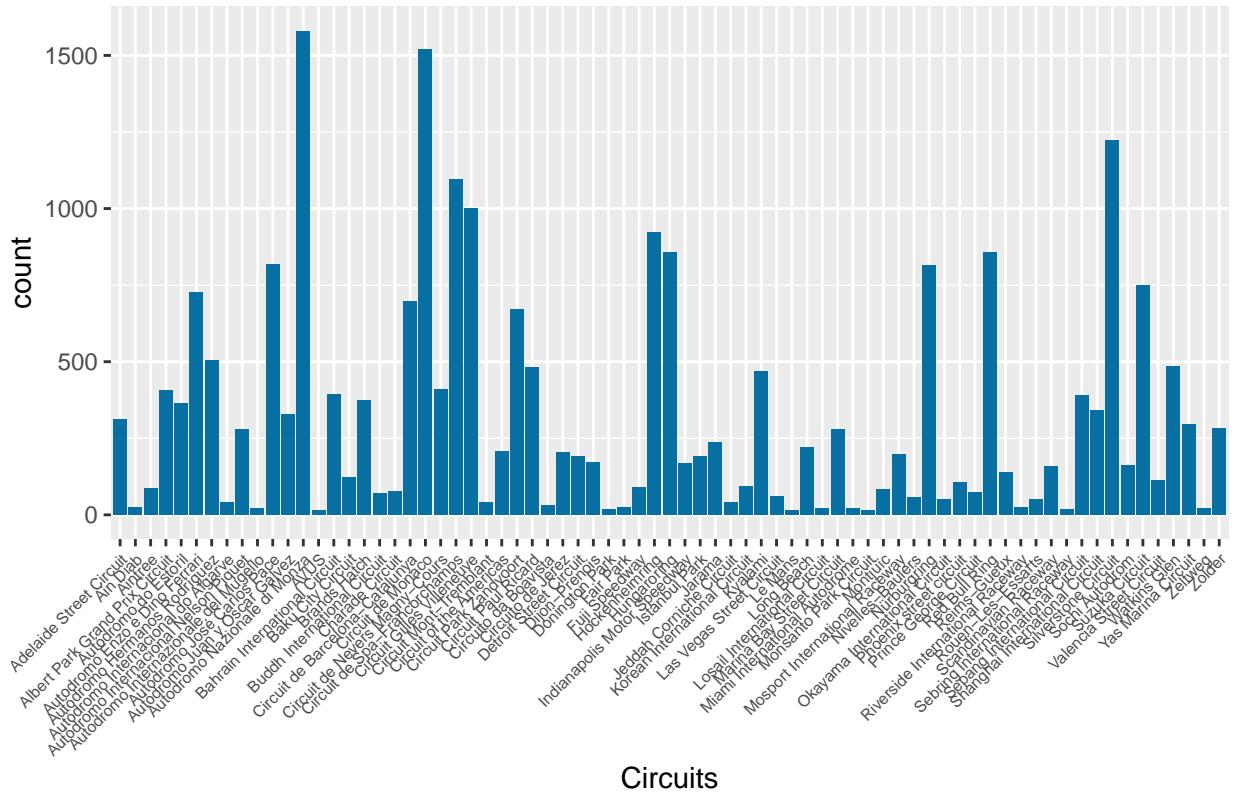
```
nationality = data.frame(table(df$drv_nationality))
nationality$Var1 = as.character(nationality$Var1)
ggplot(nationality, aes(area = Freq, fill = Freq, label = Var1)) +
  geom_treemap() +
  geom_treemap_text(fontface = "italic", colour = "white", place = "centre", grow = TRUE) +
  labs(title = "Treemap of the drivers' nationality", fill = "Frequency")
```

Treemap of the drivers' nationality



```
ggplot(df, aes(name)) +  
  labs(x = 'Circuits') +  
  ggtitle('Circuits Frequency') +  
  geom_bar(fill="#076fa2", position = position_dodge(0.7)) +  
  theme(axis.text.x = element_text(angle = 45, vjust = 1, hjust = 1, size = 6))
```

Circuits Frequency



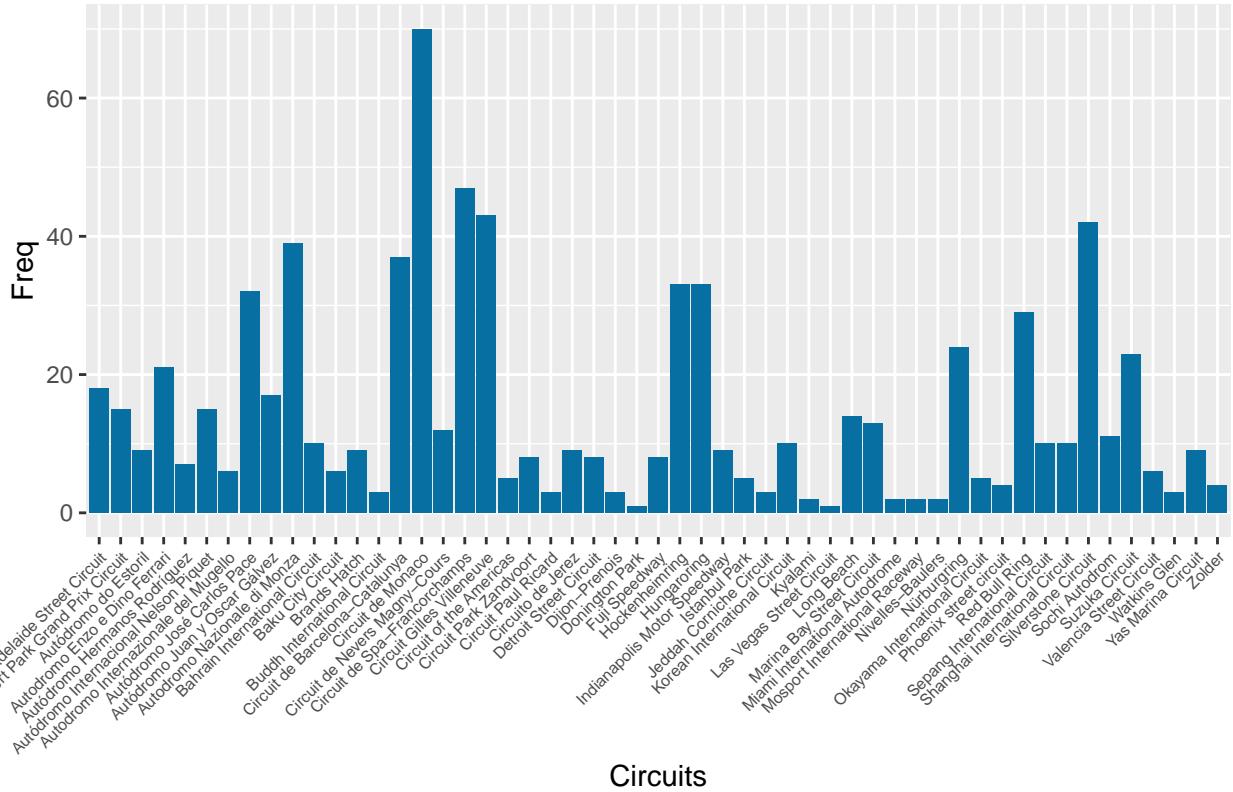
\\"

Number of collisions for each circuit.

```
#barplot showing number of collisions for each circuit
n_collision <- table(df[df$status == 'Collision', 'name'])
n_collision <- as.data.frame(n_collision)

ggplot(n_collision, aes(x = Var1, y = Freq)) +
  labs(x = 'Circuits') +
  ggtitle('Number of collisions per circuit') +
  geom_bar(fill="#076fa2", position = position_dodge(0.7), stat = 'identity') +
  theme(axis.text.x = element_text(angle = 45, vjust = 1, hjust = 1, size = 6))
```

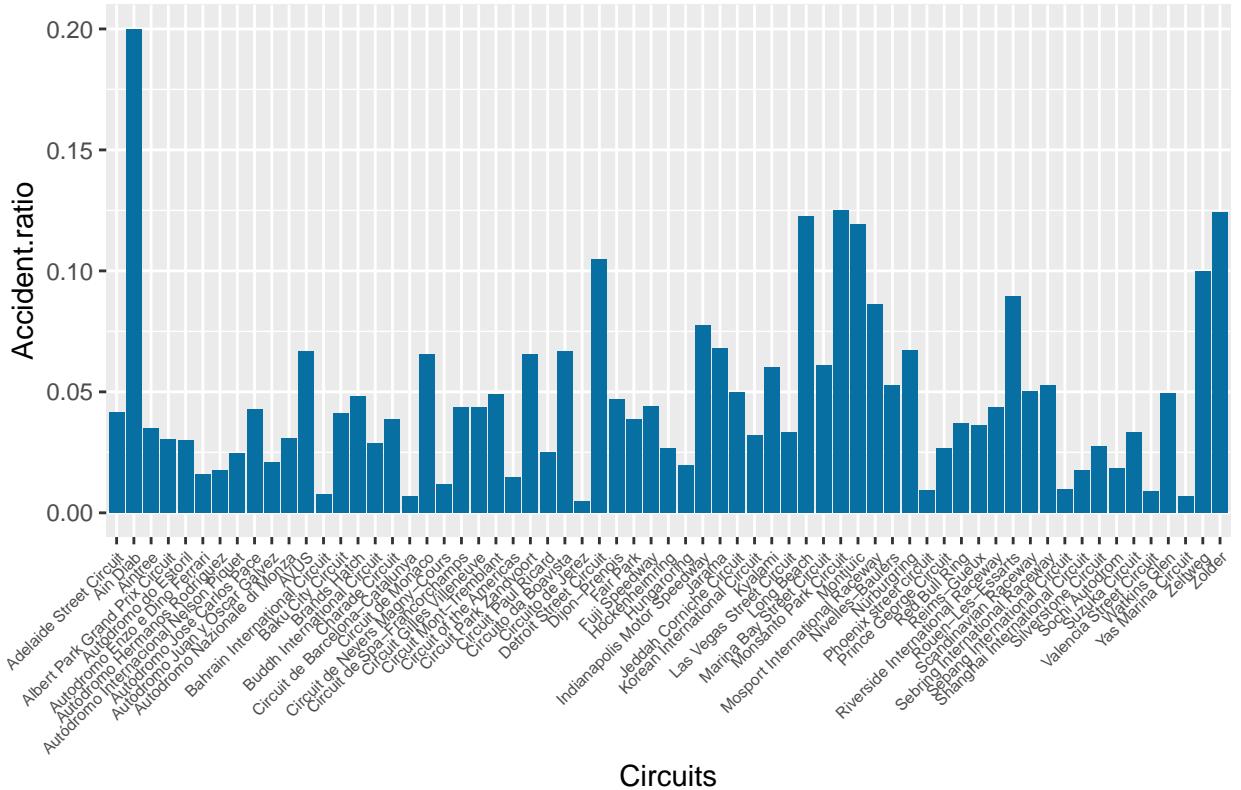
Number of collisions per circuit



Accident ratio for each circuit.

```
n_accident = read.csv("~/Library/Mobile Documents/com~apple~CloudDocs/UNICATT/Data analysis techniques/Assignment 1/n_accident.csv")
ggplot(n_accident, aes(x = as.character(Accident.Names), y = Accident.ratio)) +
  labs(x = 'Circuits') +
  ggtitle('Number of collisions per circuit') +
  geom_bar(fill="#076fa2", position = position_dodge(0.7), stat = 'identity') +
  theme(axis.text.x = element_text(angle = 45, vjust = 1, hjust = 1, size = 6))
```

Number of collisions per circuit



Winning driver age overtime

```
df_points = read.csv("~/Library/Mobile Documents/com~apple~CloudDocs/UNICATT/Data analysis techniques and applications/df_points.csv")

winner_Age = function(){
  df_year_winner = matrix(nrow = 1, ncol=5)
  years = seq(1958, 2022, 1)
  for (y in years){
    df_year = df_points[df_points$year == y, c(2,4,9,10)]
    id_winner = df_year[which.max(df_year$points), ]
    df_year_winner <- rbind(df_year_winner, c(y, id_winner$driverId, id_winner$points, id_winner$fullname))
  }
  df_year_winner = data.frame(df_year_winner)
  df_year_winner = df_year_winner[-1,]
  return(df_year_winner)
}
df_year_winner = winner_Age()
names(df_year_winner) = c('year', 'driverId', 'final_points', 'fullname', 'age')
df_year_winner$age = as.numeric(df_year_winner$age)
head(df_year_winner)

##   year driverId final_points      fullname age
## 2 1958       578            42  Mike Hawthorn 29
## 3 1959       356            31 Jack Brabham 33
```

```

## 4 1960      356      43  Jack Brabham 34
## 5 1961      403      34    Phil Hill 34
## 6 1962      289      42  Graham Hill 33
## 7 1963      373      54    Jim Clark 27

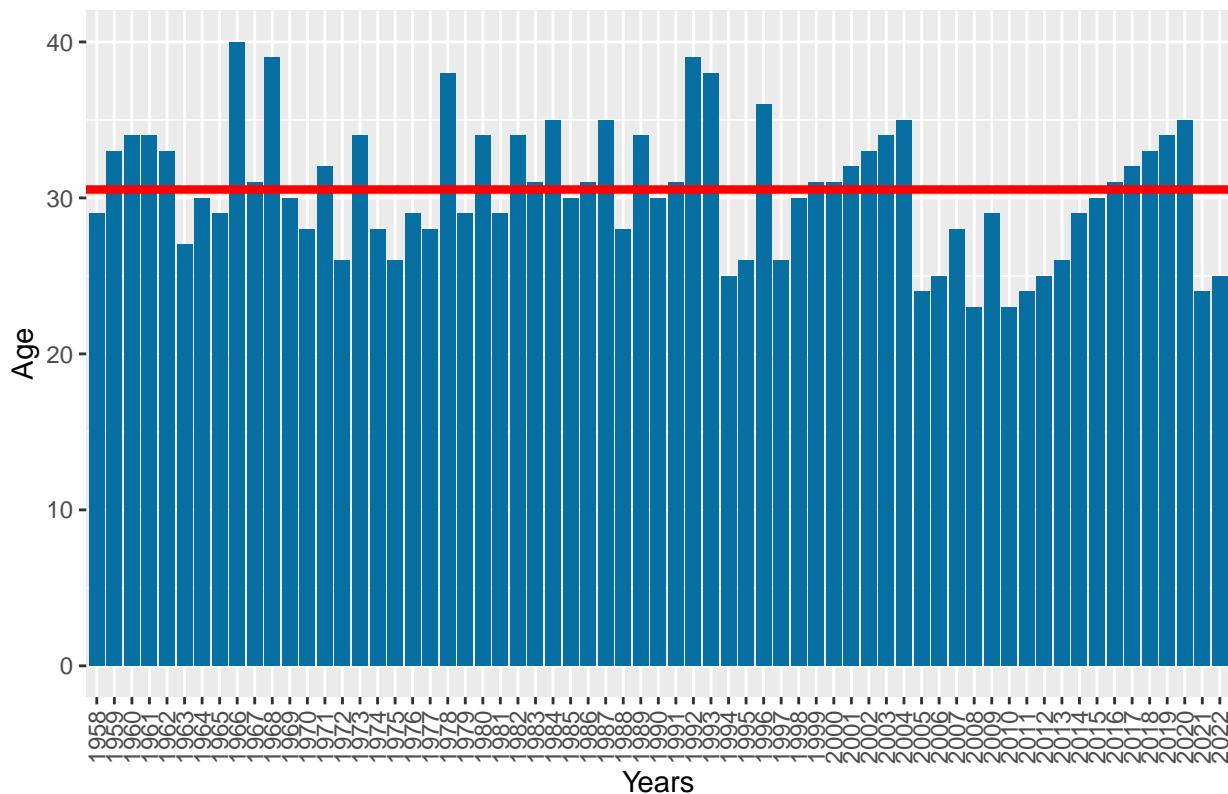
```

```

#plot for winner's age over time
ggplot(df_year_winner, aes(x = year, y = age)) +
  labs(x = 'Years', y = 'Age') +
  ggtitle('Winning driver age for every year') +
  geom_bar(stat = "identity") +
  geom_col(fill = "#076fa2") +
  theme(axis.text.x = element_text(angle = 90, vjust=0.5, hjust = 1)) +
  geom_abline(slope = 0, intercept = mean(df_year_winner$age), color = 'red', linewidth = 1.5)

```

Winning driver age for every year



How important is the pole position in each circuit to win the race?

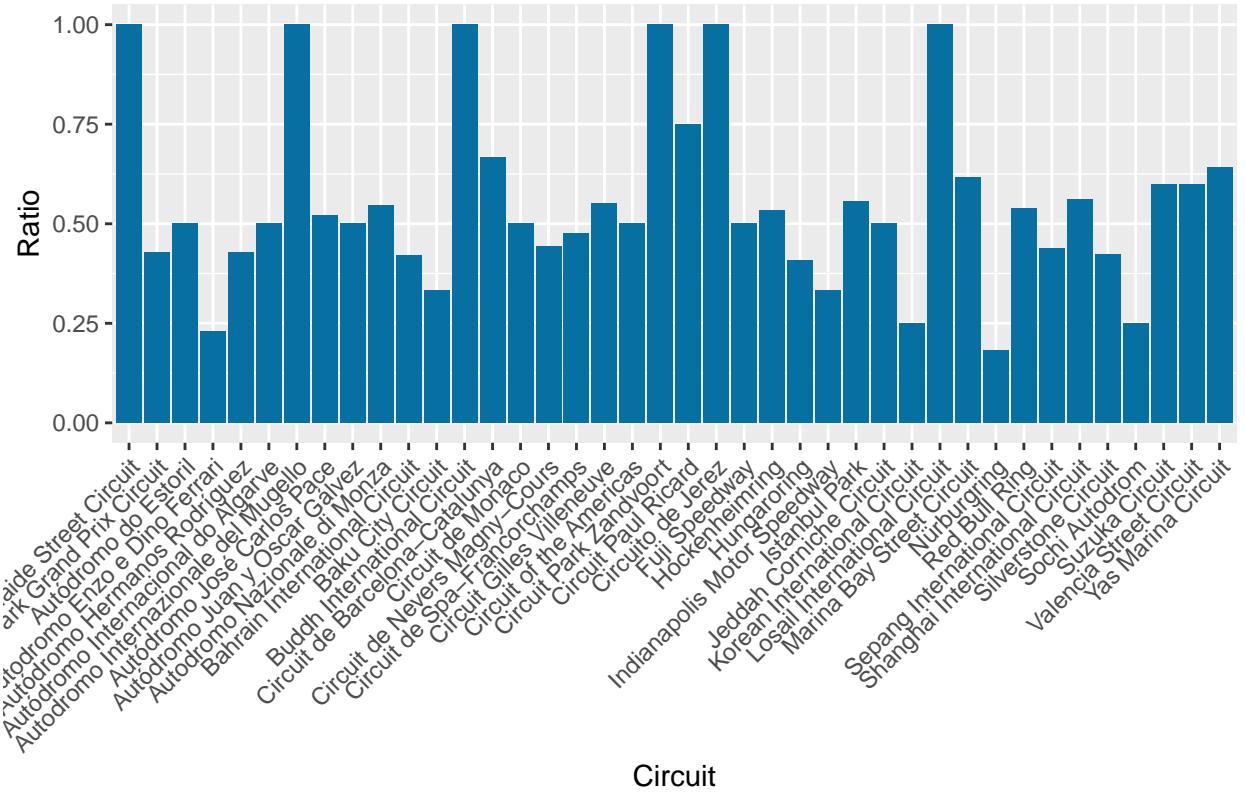
```
pole_ratio = read.csv("~/Library/Mobile Documents/com~apple~CloudDocs/UNICATT/Data analysis techniques/Formula 1 pole position.csv")
```

```

ggplot(pole_ratio, aes(x = X1, y = X2)) +
  labs(x = 'Circuit', y = 'Ratio') +
  ggtitle('How important is the pole position') +
  geom_bar(stat = "identity") +
  geom_col(fill = "#076fa2") +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))

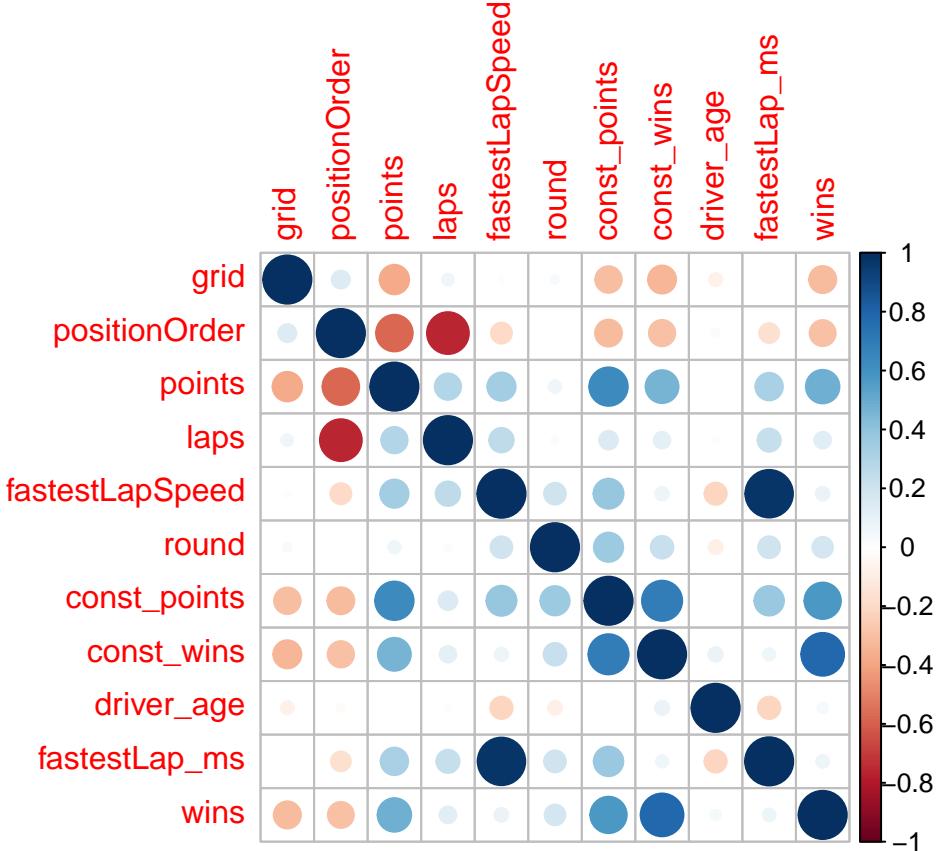
```

How important is the pole position



Correlation

```
#CORRELATION MATRIX for numerical features
corr_matrix = cor(df[c(7,8,9,10,11,16,20,21,22,23,25)])
corrplot(corr_matrix)
```



DEFINING AND TRAINING ML MODELS FOR PREDICTION.

```
df$winner = as.factor(df$winner) #converting the column to a factor.
```

Partition the data for train and validation.

For partitioning the data, we decided to use a different approach. Instead of just randomly select rows from the original dataset (we recall that each row is a driver's performance in a specific race) we decided to create a block of rows for each race (with each block containing multiple observation, i.e. all the drivers that attained the race) and randomly select 80% blocks for the training sample and the rest for the testing sample.

We decided to partition the data in this way in order to maintain the integrity of observations for each race. Otherwise, we could have situations in which in the testing sets there is just one driver for a Race, and so his performance would be assessed in the lack of opponents.

```
races_list = unique(df$raceId)
races_train_idx = sample(length(races_list), length(races_list)*0.8)

races_train_list = races_list[races_train_idx]
races_test_list = races_list[-races_train_idx]
```

```
#creating train & test dataframes
df.train = df[df$raceId %in% races_train_list, ]
df.test = df[df$raceId %in% races_test_list, ]
```

Fixing df.train and df.test for categorical features by converting those columns in factors.

```
#converting TRAIN and TEST categorical features in factors
df.train$status = as.factor(df.train$status)
df.train$driv_nationality = as.factor(df.train$driv_nationality)
df.train$fullname = as.factor(df.train$fullname)
df.train$const_name = as.factor(df.train$const_name)
df.train$name = as.factor(df.train$name)
str(df.train)
```

```
## 'data.frame': 18988 obs. of 25 variables:
##   $ raceId      : int 10 10 10 10 10 10 10 10 10 ...
##   $ driverId    : int 1 10 12 13 15 153 16 17 18 2 ...
##   $ constructorId : int 1 7 4 6 7 5 10 9 23 2 ...
##   $ circuitId   : int 11 11 11 11 11 11 11 11 11 11 ...
##   $ resultId    : int 7734 7739 7745 7753 7741 7748 7752 7736 7740 7744 ...
##   $ number       : int 1 10 8 3 9 11 20 14 22 6 ...
##   $ grid         : int 4 13 14 0 11 19 17 3 8 15 ...
##   $ positionOrder: int 1 6 12 20 8 15 19 3 7 11 ...
##   $ points       : num 10 3 0 0 1 0 0 6 2 0 ...
##   $ laps          : int 70 70 70 0 70 69 1 70 70 70 ...
##   $ fastestLapSpeed : num 191 191 189 0 189 ...
##   $ status        : Factor w/ 126 levels "+1 Lap","+10 Laps",...
##   $ dob           : chr "1985-01-07" "1982-03-18" "1985-07-25" "1981-04-25" ...
##   $ driv_nationality: Factor w/ 41 levels "American","American-Italian",...
##   $ fullname      : Factor w/ 577 levels "Adrián Campos",...
##   $ round         : int 10 10 10 10 10 10 10 10 ...
##   $ date          : chr "2009-07-26" "2009-07-26" "2009-07-26" "2009-07-26" ...
##   $ const_name    : Factor w/ 153 levels "AGS","Alfa Romeo",...
##   $ name          : Factor w/ 72 levels "Adelaide Street Circuit",...
##   $ const_points  : num 28 38.5 13 40 38.5 5 0 98.5 114 8 ...
##   $ const_wins    : int 1 0 0 0 0 0 0 3 6 0 ...
##   $ driver_age    : int 24 27 24 28 35 19 26 33 29 32 ...
##   $ fastestLap_ms : num 82479 82506 83418 0 83261 ...
##   $ winner        : Factor w/ 2 levels "0","1": 2 1 1 1 1 1 1 1 1 1 ...
##   $ wins          : int 2 1 1 1 1 1 1 2 7 1 ...
```

```
df.test$status = as.factor(df.test$status)
df.test$driv_nationality = as.factor(df.test$driv_nationality)
df.test$fullname = as.factor(df.test$fullname)
df.test$const_name = as.factor(df.test$const_name)
df.test$name = as.factor(df.test$name)
str(df.test)
```

```
## 'data.frame': 4705 obs. of 25 variables:
##   $ raceId      : int 1 1 1 1 1 1 1 1 1 ...
##   $ driverId    : int 10 15 16 17 18 2 20 21 22 3 ...
##   $ constructorId : int 7 7 10 9 23 2 9 10 23 3 ...
```

```

## $ circuitId      : int  1 1 1 1 1 1 1 1 1 1 ...
## $ resultId       : int  7557 7556 7562 7565 7554 7563 7566 7564 7555 7559 ...
## $ number         : int  10 9 20 14 22 6 15 21 23 16 ...
## $ grid           : int  19 20 16 8 1 9 3 15 2 5 ...
## $ positionOrder  : int  4 3 9 12 1 10 13 11 2 6 ...
## $ points          : num  5 6 0 0 10 0 0 0 8 3 ...
## $ laps            : int  58 58 58 57 58 58 56 58 58 58 ...
## $ fastestLapSpeed : num  216 215 215 216 217 ...
## $ status          : Factor w/ 86 levels "+1 Lap", "+11 Laps", ... : 37 37 37 1 37 37 20 37 37 37 ...
## $ dob             : chr  "1982-03-18" "1974-07-13" "1983-01-11" "1976-08-27" ...
## $ driv_nationality: Factor w/ 37 levels "American", "Argentine", ... : 16 21 16 3 7 16 16 21 6 16 ...
## $ fullname        : Factor w/ 421 levels "Adrián Campos", ... : 395 184 2 257 192 290 372 132 366 292
## $ round           : int  1 1 1 1 1 1 1 1 1 ...
## $ date            : chr  "2009-03-29" "2009-03-29" "2009-03-29" "2009-03-29" ...
## $ const_name       : Factor w/ 124 levels "AGS", "Alfa Romeo", ... : 118 118 44 99 20 12 99 44 20 122 ...
## $ name             : Factor w/ 47 levels "Adelaide Street Circuit", ... : 3 3 3 3 3 3 3 3 3 3 ...
## $ const_points     : num  11 11 0 0 18 0 0 0 18 3 ...
## $ const_wins       : int  0 0 0 0 1 0 0 0 1 0 ...
## $ driver_age       : int  27 35 26 33 29 32 22 36 37 24 ...
## $ fastestLap_ms   : num  88416 88916 88943 88508 88020 ...
## $ winner           : Factor w/ 2 levels "0", "1": 1 1 1 1 2 1 1 1 1 1 ...
## $ wins             : int  1 1 1 1 2 1 1 1 1 1 ...

```

Fixing FACTOR LEVELS in the testing dataframe.

When partitioning the data, categories for nominal features are shuffled between df.train and df.test. This lead to a situation in which categories in df.test could not be present in df.train and so classification models trained on df.test would not consider other categories that are only present in df.test.

With the following lines of code we overcome this problem by setting factor levels of all categorical columns of df.test to be exactly the same as df.train. Otherwise, we would encounter problems in the classifications method when predicting the new data from df.test using the model trained on df.train.

```

levels(df.test$status) <- levels(df.train$status)
levels(df.test$driv_nationality) <- levels(df.train$driv_nationality)
levels(df.test$fullname) <- levels(df.train$fullname)
levels(df.test$const_name) <- levels(df.train$const_name)
levels(df.test$name) <- levels(df.train$name)

```

SCORING FUNCTIONS

The two scoring functions we created, simply calculate the accuracy for the prediction of the winner of each race.

For the regressions, the scoring functions sort the prediction of `positionOrder` for each race (by using `raceId`) and take the smallest value for the prediction (since in the `positionsOrder` the winner has position 1, hence the smallest position). If the same driver (a row of df.test) has the real value of `positionOrder` equal to 1, then the prediction is correct and the score is increased by 1. After scrolling through all races that are contained in df.test, the `accuracy_score` is defined as the ratio between score and the number of races.

For the classification, the scoring function works similarly but there is a difference. We are using the classifications models to predict `winner` defined in {0,1} but the models in action are going to assign 1 to multiple drivers in the same race (but we can have just one winner in a race). So instead we look for the greatest probability among the driver of a race to have the `winner` feature equal to 1. The driver with the

greatest probability of having the `winner` feature equal to 1 would be the driver with the greatest probability of winning the race and hence he is the winner our function pick. The definition of score works in the same way as in the scoring function for regressions.

Defining scoring function for Regression.

```
score_regression <- function(test, prediction){

  df.score.regres = data.frame(test, prediction)
  colnames(df.score.regres)[c(26)] = c('prediction')

  racelist = unique(df.score.regres$raceId)

  score = 0
  for (i in 1:length(racelist)){
    race = df.score.regres[df.score.regres$raceId %in% racelist[i], ]
    pred_winner_idx = which.min(race$prediction)
    if (race$positionOrder[pred_winner_idx] == 1){
      score = score + 1
    }
  }

  score_ratio = score/length(racelist)
  return(score_ratio)
}
```

Defining scoring function for Classification.

```
score_classification <- function(test, prediction){

  df.score.class = data.frame(test, prediction)
  colnames(df.score.class)[c(26,27)] = c('pred_0', 'pred_1')

  racelist = unique(df.score.class$raceId)

  score = 0
  for (i in 1:length(racelist)){
    race = df.score.class[df.score.class$raceId %in% racelist[i], ]
    pred_winner_idx = which.max(race$pred_1)
    if (race$positionOrder[pred_winner_idx] == 1){
      score = score + 1
    }
  }

  score_ratio = score/length(racelist)
  return(score_ratio)
}
```

REGRESSION

Training and testing linear regression models. For the regression model we decided to use as explanatory variables most of the numerical features as: grid, laps, fastestLapSpeed, round, const_points, const_wins, fastestLap_ms and wins that presumably would help predict and explain the final position of

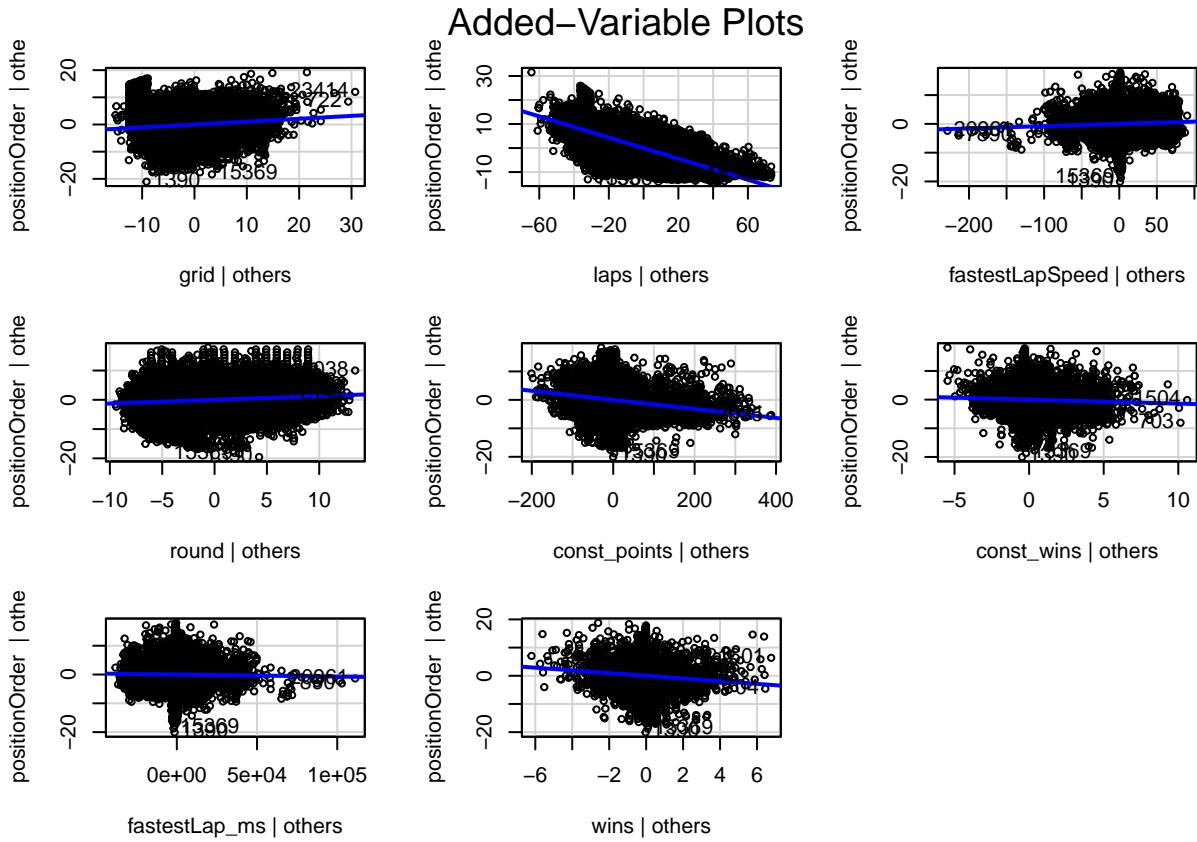
the driver. Some variables were excluded as points, since its high correlation with the final position (higher number of points means a higher position, where 1 is the highest position).

```
#first, we train and test a linear model with all numerical features in the dataset.  
#linear regression
```

```
linearmodel = lm(positionOrder ~ grid + laps + fastestLapSpeed + round + const_points + const_wins + fa  
summary(linearmodel)
```

```
##  
## Call:  
## lm(formula = positionOrder ~ grid + laps + fastestLapSpeed +  
##       round + const_points + const_wins + fastestLap_ms + wins,  
##       data = df.train)  
##  
## Residuals:  
##      Min        1Q    Median        3Q       Max  
## -20.0893  -2.9626  -0.0301   2.9269  17.9105  
##  
## Coefficients:  
##              Estimate Std. Error t value Pr(>|t|)  
## (Intercept) 2.137e+01 1.119e-01 190.971 < 2e-16 ***  
## grid        1.055e-01 5.115e-03 20.631 < 2e-16 ***  
## laps        -2.198e-01 1.360e-03 -161.585 < 2e-16 ***  
## fastestLapSpeed 7.931e-03 1.654e-03  4.794 1.65e-06 ***  
## round       1.268e-01 7.411e-03 17.108 < 2e-16 ***  
## const_points -1.590e-02 7.028e-04 -22.627 < 2e-16 ***  
## const_wins   -1.395e-01 3.354e-02 -4.160 3.20e-05 ***  
## fastestLap_ms -7.174e-06 3.616e-06 -1.984 0.0473 *  
## wins         -4.777e-01 4.782e-02 -9.991 < 2e-16 ***  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
##  
## Residual standard error: 4.6 on 18979 degrees of freedom  
## Multiple R-squared:  0.643, Adjusted R-squared:  0.6428  
## F-statistic: 4272 on 8 and 18979 DF, p-value: < 2.2e-16
```

```
avPlots(linearmodel)
```



As we can see the linear model obtain an R-squared of 65% that confirms a good capacity of the model to represent variability in the original data. Also, the majority of variables are strongly significant. We can now try to run a prediction and recall ‘score_regression’ function to see the results of the prediction.

```
predict_lm = predict(linearmodel, df.test)
score_regression(df.test, predict_lm)
```

```
## [1] 0.6600985
```

Regression tree Decision trees are another technique that we have used both in regression and in classification. The model used in regression take as explanatory the same variables in the regression:

```
predict_rt = predict(tree(positionOrder ~ grid + number + laps + fastestLapSpeed + round + const_points
                           , df.train), newdata = df.test)
score_regression(df.test, predict_rt) #47%
```

```
## [1] 0.4384236
```

CLASSIFICATION

Using the features class **winner** that we have created, we want to predict the probability of class=1 (winner) or class=0 (not winner). Then we are going to sort the probabilities and pick the greater probability of class=1, hence the driver with the greatest probability of being a winner.

For the classification methods, we decided to remove those categorical features as `status`, `fullname` and `const_name` that when converted in factors were causing problems in the application of the techniques, or drastically reduced performances due to their dimension (all of those 3 factors presented more than 100 levels that were not accepted by the functions or ruined the prediction accuracy).

```
df.nb <- naiveBayes(winner ~ . -number -positionOrder -resultId -points -status -fullname -const_name, df)
prediction_nb <- predict(df.nb, newdata = df.test, type = 'raw')
score_classification(df.test, prediction_nb)
```

BAYES CLASSIFIER

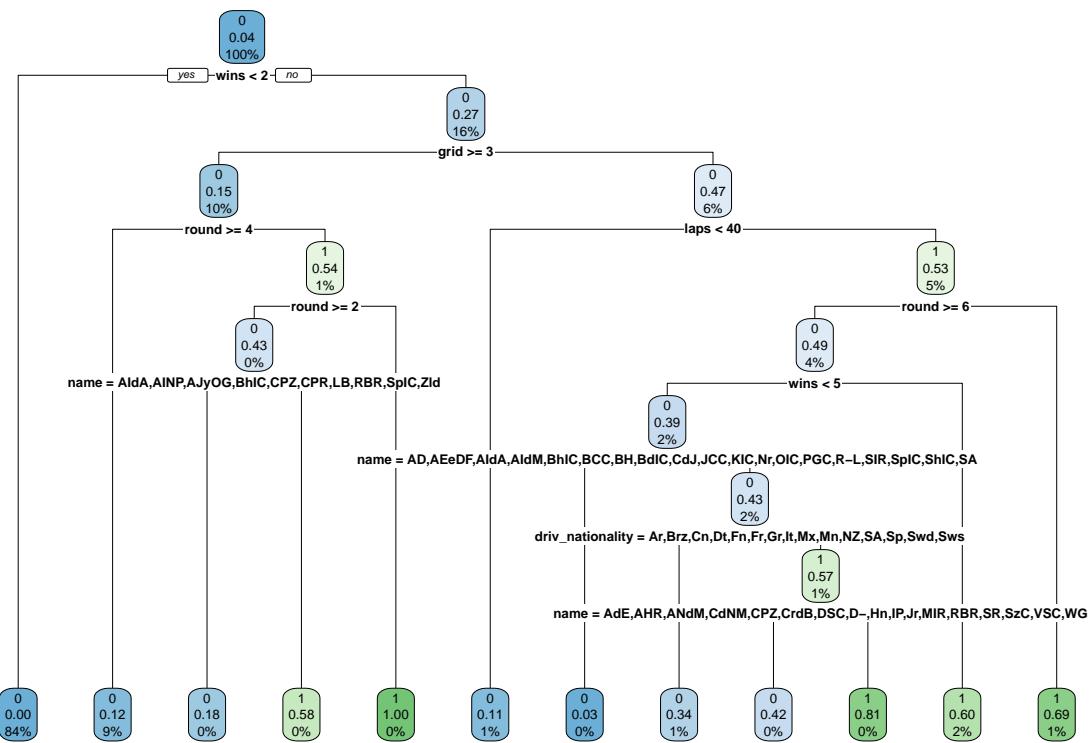
```
## [1] 0.5172414
```

```
df.dt = rpart(winner ~ . -number -positionOrder -resultId -points -status -fullname -const_name -dob -date, df)
prediction_dt = predict(df.dt, newdata = df.test, method="prob")
score_classification(df.test, prediction_dt) #58%
```

DECISION TREES

```
## [1] 0.591133
```

```
rpart.plot(df.dt,facelen = 2)
```



```
df.rf <- randomForest(winner ~ . -number -points -positionOrder -resultId -const_name -name -fullname -  
prediction.rf <- predict(df.rf, df.test, type = 'prob')  
prediction.rf[is.na(prediction.rf)] <- 0  
score_classification(df.test, prediction.rf)
```

RANDOM FOREST

```
## [1] 0.6699507
```

```
df.lsvm = svm(winner ~ ., data = df.train[, -c(5,6,8,9,12,13,15,17)], kernel = 'linear', fitted = FALSE  
prediction_svm <- predict(df.lsvm, newdata = df.test[,-c(5,6,8,9,12,13,15,17)], fitted = FALSE, probability = TRUE)  
SVM_class = data.frame(attributes(prediction_svm)$probabilities)  
score_classification(df.test, SVM_class)
```

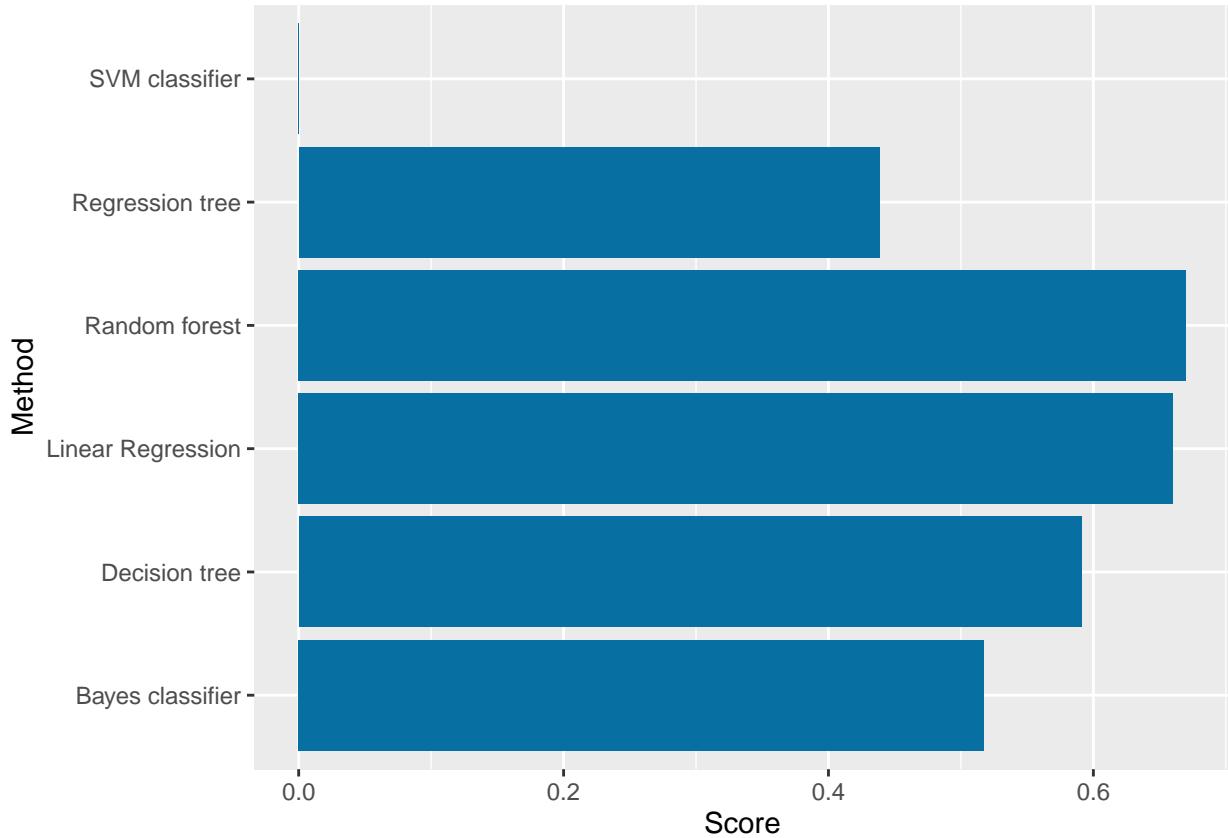
SVM classifier

```
## [1] 0
```

Visualize final prediction results

```
accuracy_results = data.frame(c(score_regression(df.test, predict_lm), score_regression(df.test, predict_rf),
colnames(accuracy_results) = c('score', 'method')

ggplot(accuracy_results, aes(x = score, y = method)) +
  labs(x = 'Score', y = 'Method') +
  geom_bar(stat = "identity") +
  geom_col(fill = "#076fa2")
```



CONCLUSIONS

Our initial goal wasn't to analyze deeply the data that characterize a Formula 1 Race. Our goal was to rather try to predict the winner of a Race with the least possible number of numerical and categorical features.